



Predicting Flight Delays

[W261 Final Project Link](#)

Team 28: Aditya Bajaj, Gerrit Lensink, Ruby Han, Tsun Tsai

Outline

01

Project Description

03

**Performance and Cap
Analysis**

02

Pipeline

04

**Limitations,
Challenges & Future
Work**

Project Description



\$30 billion

Industry losses as a result of flight delays



Delay reduction

Motivation for airlines to reduce flight delays

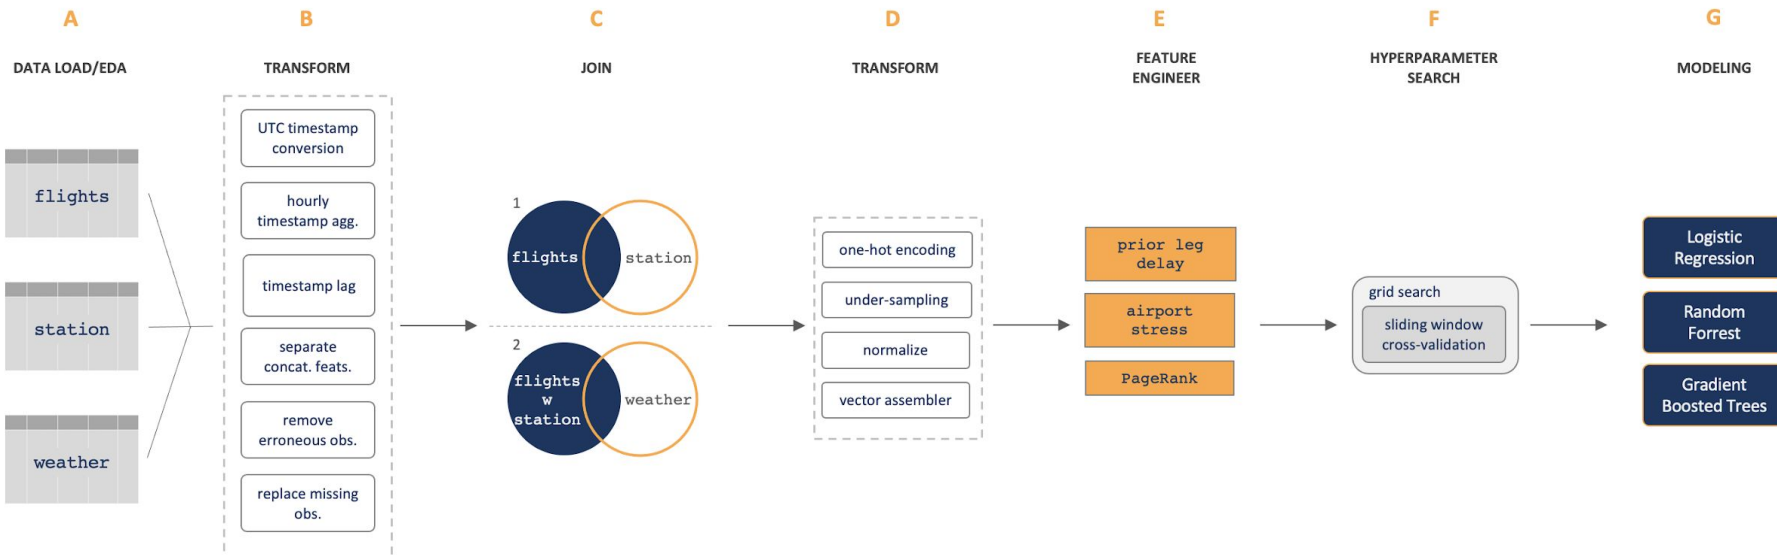


Airlines

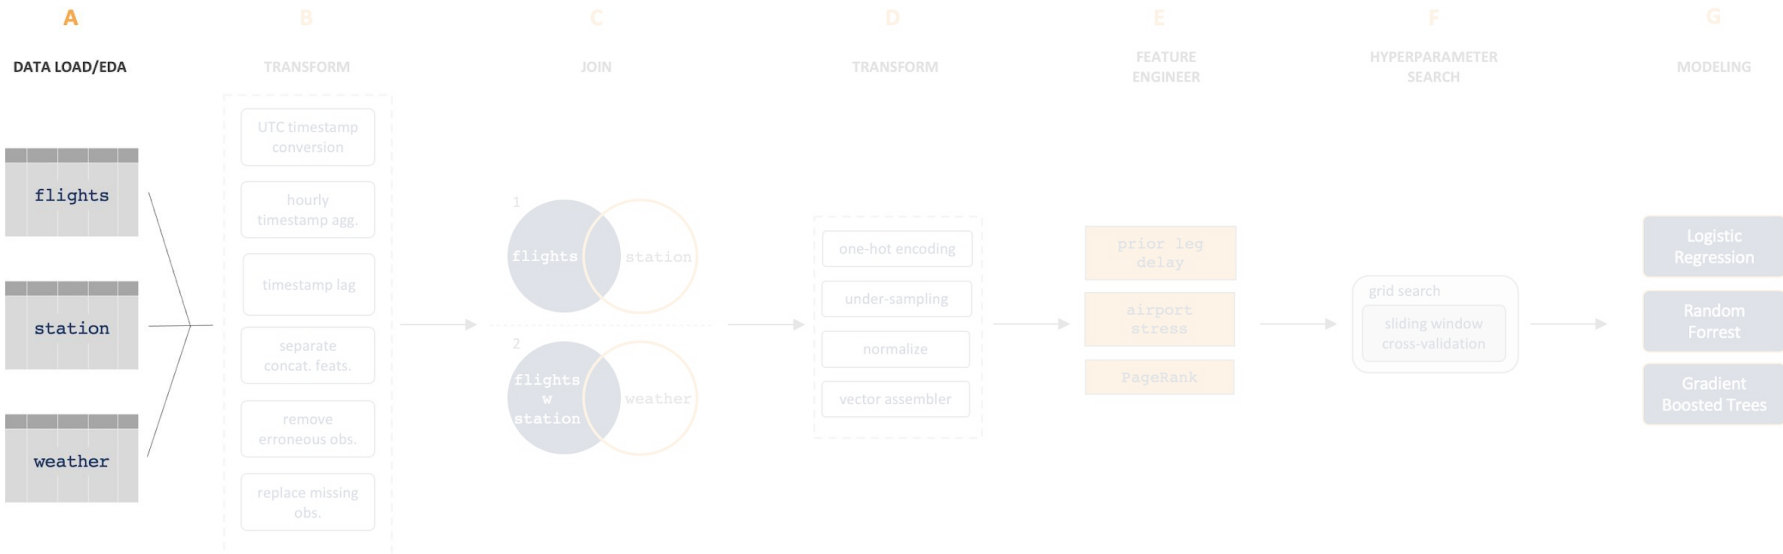
Create a model for stakeholder to predict flight delays, when unavoidable



Pipeline




Pipeline - Data Load/EDA

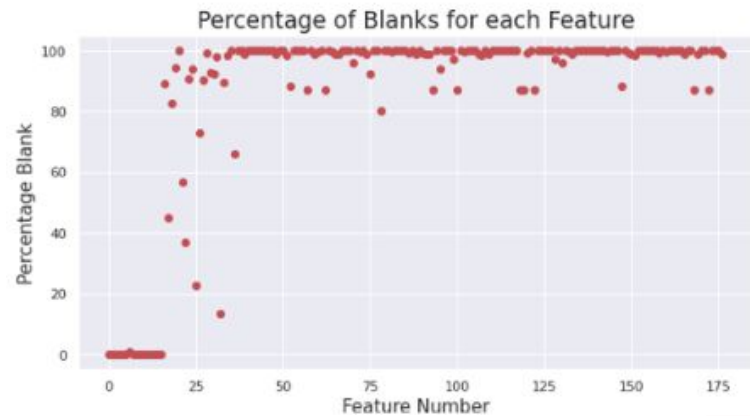


Pipeline - Data Load/EDA


CLASS IMBALANCE




NULL DATA



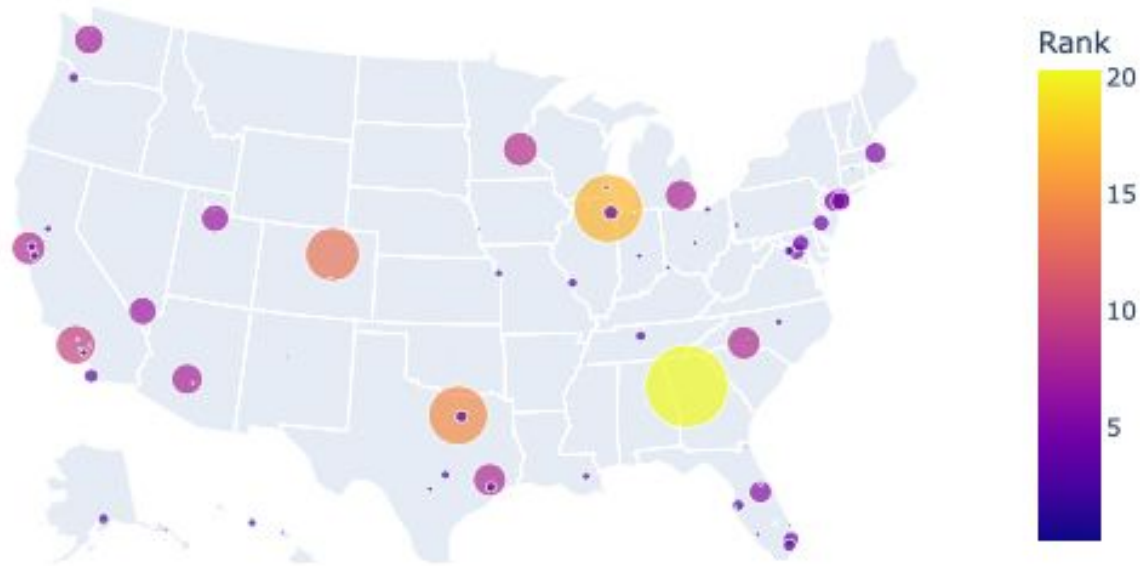
Pipeline - Data Load/EDA



Pipeline - Data Load/EDA



Pipeline - Data Load/EDA



Pipeline - Transform I



Pipeline - Transform I

Case	Strategy	Airline			"Hourly" UTC Time (Lagged)	Weather Worst Case UTC Time	Difference in Actual UTC Times
		UTC Time	"Hourly" UTC Time	Hour Lag			
1	Round-down	1:59 PM	1:00 PM	2	11:00 AM	11:59 AM	2 hr 0 min
		1:01 PM					1hr 2 min
2	Round-down	1:59 PM	1:00 PM	3	10:00 AM	10:59 AM	3 hr 0 min
		1:01 PM					2 hr 2 min



Pipeline - Transform I + II



TRANSFORMATION

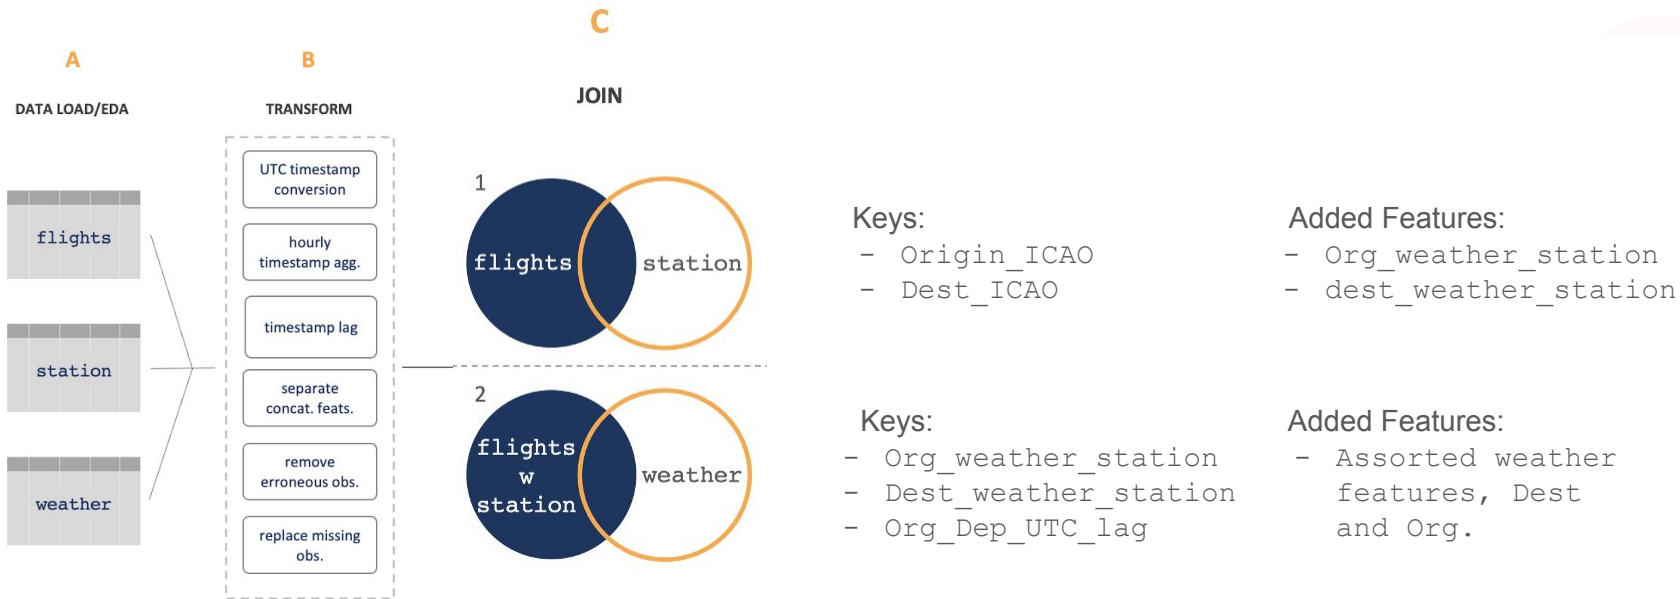
- UTC timestamp conversion
- **Hourly timestamp aggregation**
- **Timestamp lag**
- Cross-map to fill `station_id` missing values
- Parse concatenated features
- Remove erroneous observations
- **Replace missing observations**
 - *Seven-day look back*
- **One-hot encoding**
- **Undersampling**
- Scaling continuous features



ENGINEERED FEATURES

- Prior leg departure/arrival delay
- Airport stress @ departure/arrival time
- PageRank (airport traffic score)
- Dew & air temperature difference
- Vectorized feature list


Pipeline - Join



Data - Join datasets

JOIN 1:

flights_w_station



Keys:

- Origin_ICAO
- Dest_ICAO

Added Features:

- Org_weather_station
- dest_weather_station

JOIN 2:

flights_w_weather



Keys:

- Org_weather_station
- Dest_weather_station
- Org_Dep_UTC_lag

Added Features:

- Assorted weather features, Dest and Org.

Data - Join datasets

```
# Second, we need to create a new timestamp to join weather on - which is lagged by 2 hours.
# Due to the nature of our timestamps - we will actually have to subtract 3 hours to avoid leakage. 10800 is equivalent to 3 hours
flights_w_stations = flights_w_stations.withColumn('CRS_DEP_TIME_UTC_LAG', to_timestamp(f.col('CRS_DEP_TIME_UTC_HOUR').cast('long')
- 10800))

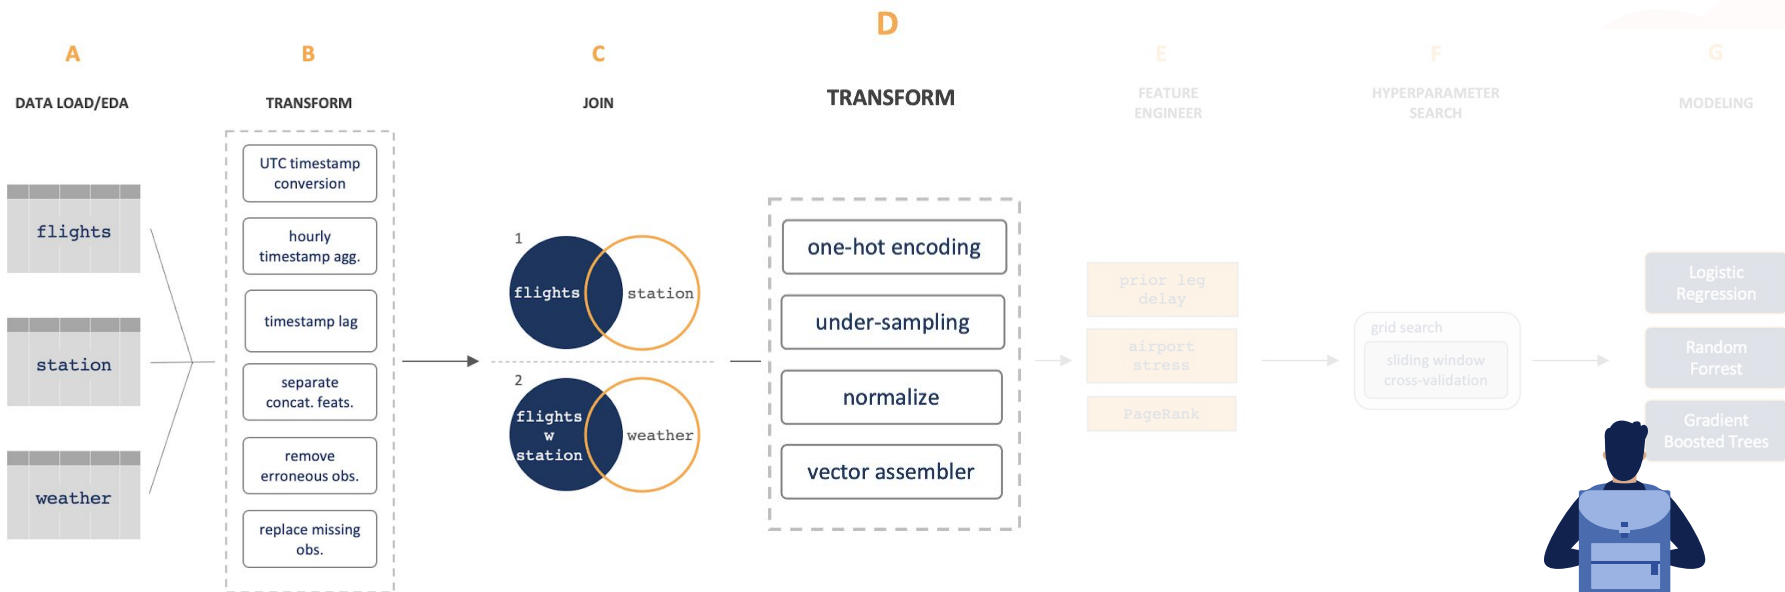
# Next, prepend an origin or destination prefix to the weather columns, so when we join we know which weather set we're looking at
origin_weather = weather.select([f.col(weather_feat).alias('ORIGIN_WEATHER_'+weather_feat) for weather_feat in weather.columns])
dest_weather = weather.select([f.col(weather_feat).alias('DEST_WEATHER_'+weather_feat) for weather_feat in weather.columns])

# join flights to ORIGIN weather on station_id and our lagged ORIGIN time variable
flights_w_weather_temp = flights_w_stations.join(origin_weather, (flights_w_stations.ORIGIN_WEATHER_STATION_ID ==
origin_weather.ORIGIN_WEATHER_STATION) & (flights_w_stations.CRS_DEP_TIME_UTC_LAG == origin_weather.ORIGIN_WEATHER_HOUR), 'left')

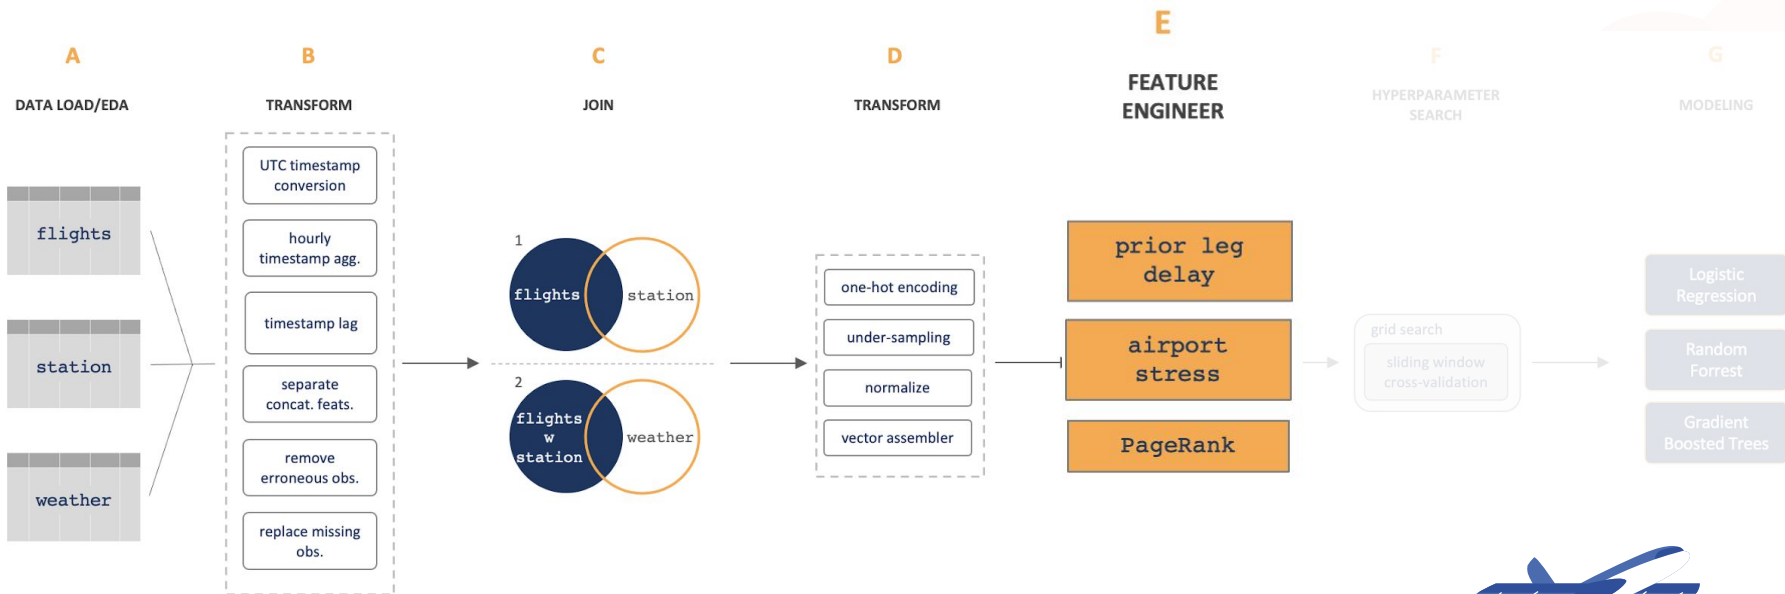
# Finally, join flights to DESTINATION weather on station_id and lagged ORIGIN time variable
flights_w_weather = flights_w_weather_temp.join(dest_weather, (flights_w_weather_temp.DEST_WEATHER_STATION_ID ==
dest_weather.DEST_WEATHER_STATION) & (flights_w_weather_temp.CRS_DEP_TIME_UTC_LAG == dest_weather.DEST_WEATHER_HOUR), 'left')
```



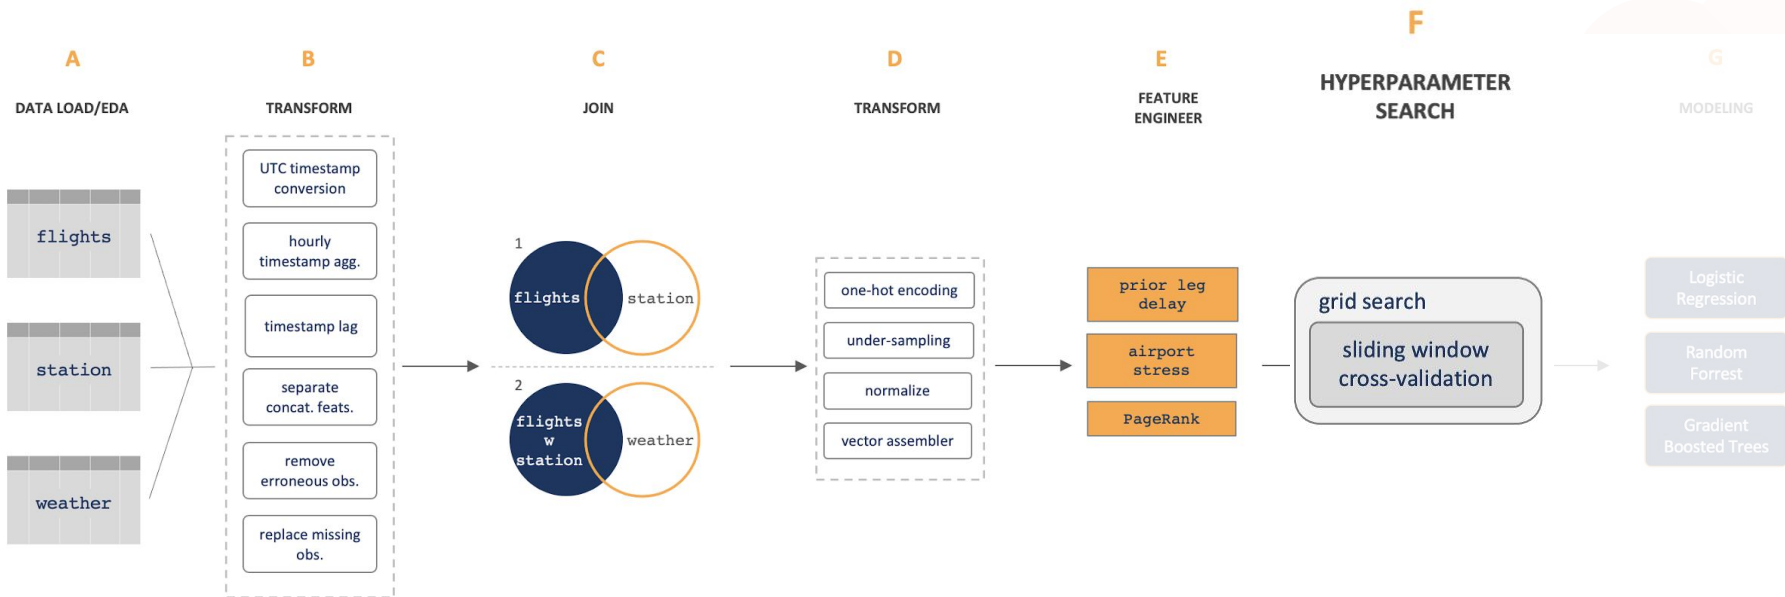
Pipeline - Transform II



Pipeline - Feature Engineering




Pipeline - Hyperparameter Search



Pipeline - Hyperparameter Search

5-fold Sliding Window CV





Model

Model - CV and Grid Search

```
def cv_folds(self):  
  
    # Create simple folds by full train year, val year  
    if self.simple_splits is True:  
        for year in [2015, 2016, 2017]:  
  
            train_df = self.df.where(f.col('CRS_DEP_TIME_UTC_HOUR') == year)  
            val_df = self.df.where(f.col('CRS_DEP_TIME_UTC_HOUR') == (year + 1))  
  
            self.train_df.append(train_df)  
            self.test_df.append(val_df)  
  
    # Create folds by number of folds, train/test split %  
    else:  
        data_size = self.df.count()  
  
        for k in range(self.folds):  
  
            # Calculate total size of each fold  
            fold_size = data_size/self.folds  
  
            # Find range of 'SPLIT_ID' for train and val  
            train_ids = ((fold_size * k) + 0,  
                        int(fold_size * k + fold_size*self.train_percent))  
  
            val_ids = (train_ids[1] + 1,  
                      fold_size * k + fold_size)  
  
            # Split the data  
            train_df = self.df.where(f.col('SPLIT_ID').between(train_ids[0], train_ids[1]))  
            val_df = self.df.where(f.col('SPLIT_ID').between(val_ids[0], val_ids[1]))  
  
            # store data  
            train_df = DataSplit.class_balance(train_df, 'undersample')  
            self.train_df.append(train_df)  
            self.test_df.append(val_df)  
  
    return train_dfs, val_dfs
```



Model - CV and Grid Search

```
def gridsearch(full_data, k_folds, param_list, param_names, random_shuffle_top_N=10, threshold_plot=False, model_type='lr'):
```

```
    # Assemble train/test from cv_folds
```

```
    model = Model(df=full_data, folds = k_folds, train_percent = .8, simple_splits = False)
    model.cv_folds()
```

```
    # Assemble Gridsearch
```

```
    params = list(itertools.product(*param_list))
    random.shuffle(params)
```

```
    # Run Models in folds
```

```
    for param in params[:random_shuffle_top_N]:
```

```
        print(f"\nLooping through param:{list(zip(param_names,param))}")
```

```
        if model_type == 'lr':
```

```
            model.get_logistic_regression(**{'maxIter':param[0], 'elasticNetParam':param[1], 'regParam':param[2]})
```

```
        if model_type == 'rf':
```

```
            model.get_random_forest(**{'numTrees':param[0], 'maxDepth':param[1], 'maxBins':param[2]})
```

```
        if model_type == 'gbt':
```

```
            model.get_GBT(**{'maxIter':param[0], 'maxDepth':param[1], 'maxBins':param[2], 'stepSize':param[3]})
```

```
    print(f'GridSearch\nBest Valid f0.5 Score: {model.best_score:.3f}\nBest Parameters: {model.best_params}')
```

```
    if threshold_plot:
```

```
        preds_valid_PR = model.tune_threshold()
```

```
        model.threshold_plot(preds_valid_PR)
```



Model - CV and Grid Search

```
class Model(DataSplit):
```

```
# def __init__(self, train_df, test_df, label_col=['DEP_DEL15'], feature_col=['ALL_FEATURES_VA'],
**kwargs):
    def __init__(self, df, label_col=['DEP_DEL15'], feature_col=['ALL_FEATURES_VA'], train_percent=0.8,
timestamp_col='CRS_DEP_TIME_UTC', folds = 1, simple_splits = False):
```

```
        super().__init__(df, train_percent, timestamp_col, folds, simple_splits)
        self.label_col = label_col
        self.feature_col = feature_col
        self.model = None
        self.best_score = 0
        self.best_params = None
        self.best_preds_test = None
```

```
def get_logistic_regression(self, **kwargs):
```

```
    model_args = kwargs
    lr = LogisticRegression(label_col = self.label_col[0],
                            features_col=self.feature_col[0],
                            elasticNetParam = model_args.get('elasticNetParam',1),
                            standardization = model_args.get('standardization',True),
                            maxIter=model_args.get('maxIter',10),
                            regParam=model_args.get('regParam',0.001))
```

```
    for index in range(len(self.train_df)):
        print(f'\n{index}-Fold:\n')
        self.model = lr.fit(self.train_df[index])
        self.get_predictions(index, model_args)
```

```
def get_random_forest(self, **kwargs):
```

```
    model_args = kwargs
    rf = RandomForestClassifier(label_col = self.label_col[0],
                               features_col=self.feature_col[0],
                               numTrees=model_args.get('numTrees',1),
                               maxDepth=model_args.get('maxDepth',5),
                               maxBins=model_args.get('maxBins',32))
```

```
    for index in range(len(self.train_df)):
        print(f'\n{index}-Fold:\n')
        self.model = rf.fit(self.train_df[index])
        self.get_predictions(index, model_args)
```

```
def get_GBT(self, **kwargs):
```

```
    model_args = kwargs
    GBT = GBTClassifier(label_col = self.label_col[0],
                        features_col=self.feature_col[0],
                        maxIter=model_args.get('maxIter',1),
                        maxDepth=model_args.get('maxDepth',5),
                        maxBins=model_args.get('maxBins',32),
                        stepSize=model_args.get('stepSize',0.1))
```

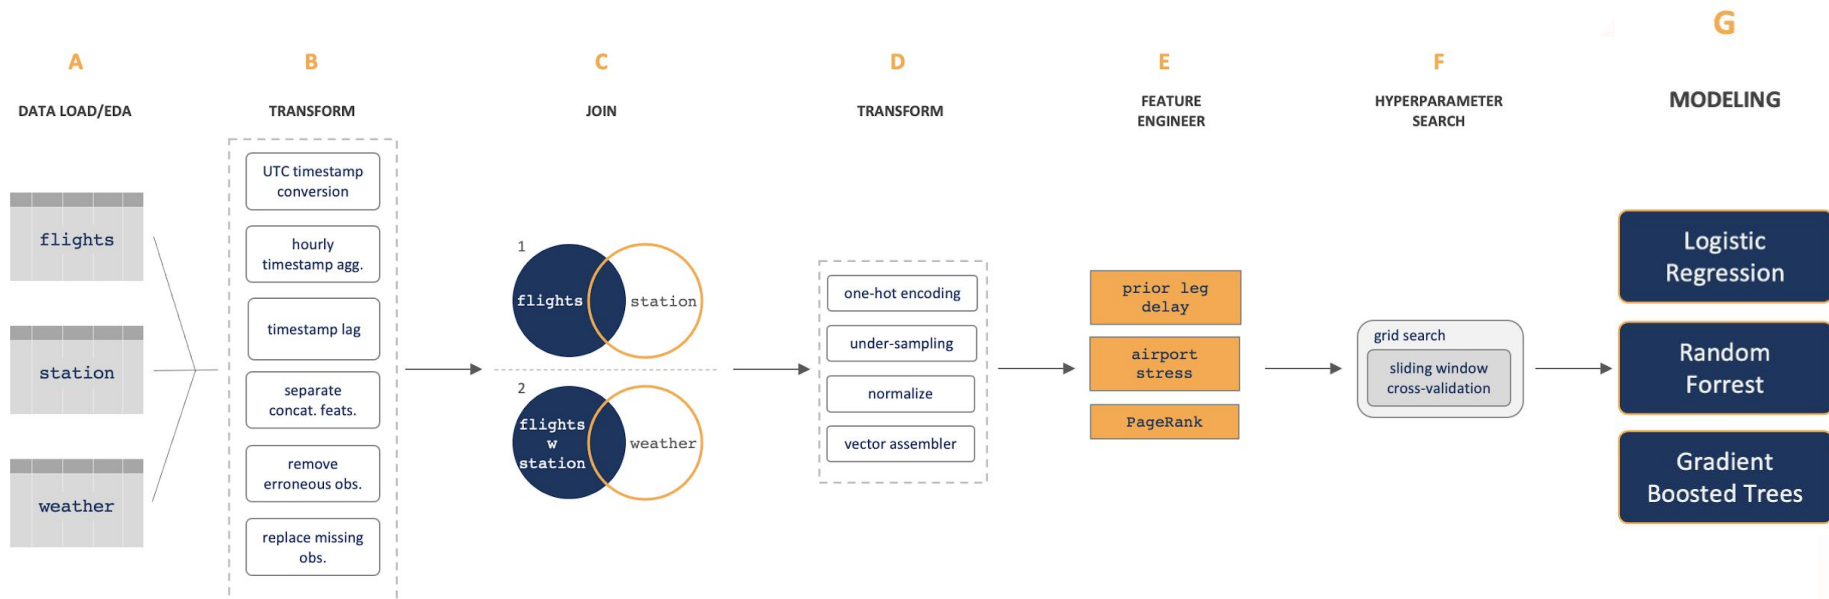
```
    for index in range(len(self.train_df)):
        print(f'\n{index}-Fold:\n')
        self.model = GBT.fit(self.train_df[index])
        self.get_predictions(index, model_args)
```

```
def get_metrics(self, preds_train, preds_test, model_args):
```

```
    train_rdd = preds_train.select(['prediction', 'DEP_DEL15']).rdd
    train_metrics = MulticlassMetrics(train_rdd)
    valid_rdd = preds_test.select(['prediction', 'DEP_DEL15']).rdd
    valid_metrics = MulticlassMetrics(valid_rdd)
    valid_score = valid_metrics.fMeasure(1.0, 0.5)
    if valid_score > self.best_score:
        self.best_score = valid_score
        self.best_params = model_args
        self.best_preds_test = preds_test
        print(' \t\tTrain Metrics \t Validation Metrics\n')
        print(f'Recall: \t\t{train_metrics.recall(label=1):.3f} \t\t{valid_metrics.recall(label=1):.3f}')
        print(f'Precision: \t\t{train_metrics.precision(1):.3f} \t\t{valid_metrics.precision(1):.3f}')
        print(f'Accuracy: \t\t{train_metrics.accuracy:.3f} \t\t{valid_metrics.accuracy:.3f}')
        print(f'F0.5 score: \t\t{train_metrics.fMeasure(1.0, 0.5):.3f} \t\t{valid_metrics.fMeasure(1.0, 0.5):.3f}')
        print(f'F2 score: \t\t{train_metrics.fMeasure(1.0, 2.0):.3f} \t\t{valid_metrics.fMeasure(1.0, 2.0):.3f}')
        print(f'F1 score: \t\t{train_metrics.fMeasure(1.0):.3f} \t\t{valid_metrics.fMeasure(1.0):.3f}')
```



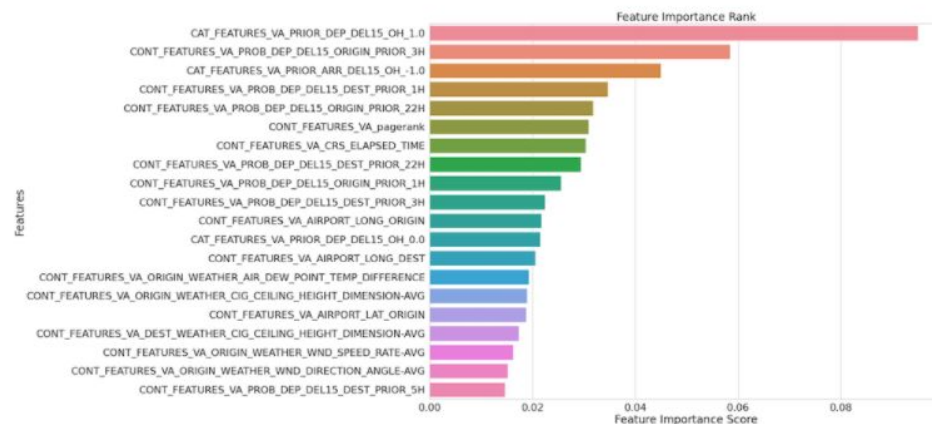
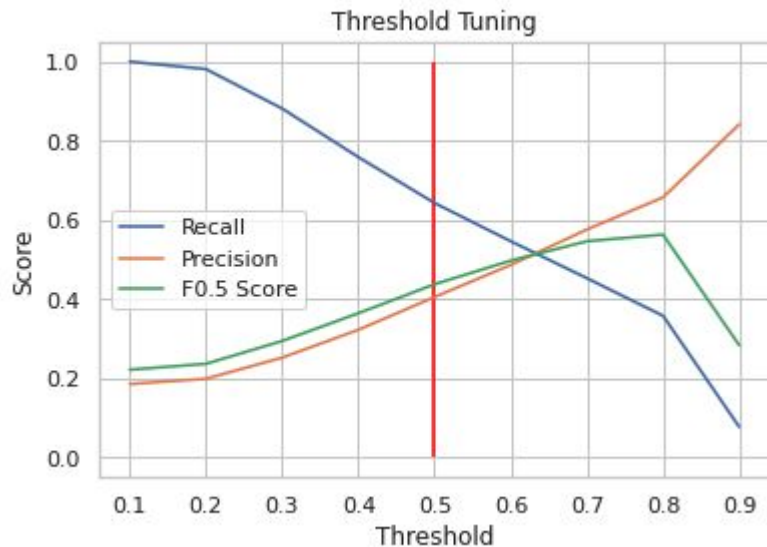
Pipeline - Modeling





Model Results

Pipeline - Modeling



Pipeline - Modeling

Phase	Model #	Algorithm	F0.5	CV Folds	GridSearch	Param Combos	RunTime	Change log	Train	Val	Test
II	Initial	LR	0.049	1	N	-	-	Did not account for class imbalance, regularization, OHE, VA	Jan - Sep 2015	Oct - Dec 2015	N/A
	Baseline	LR	0.253	1	N	-	-	Add undersampling (~20% sampled on on-time flights)			
	1	LR	0.453*	1	N	-	-	Add regularization, OHE, VA			
	2a	LR	0.208*	1	N	-	-	Add undersampling			
	2b	RF	0.244	1	N	-	-				
III	3	LR	0.284	1	N	-	-	Change Train/Val	2017	2018	N/A
	4a	LR	0.304	1	N	-	-	Add PRIOR_ARR_DEL15			
	4b	RF	0.281	1	N	-	-				
	5a	LR	0.384	1	N	-	-	Add PRIOR_DEP_DEL15			
	5b	RF	0.362	1	N	-	-				
	5c	GBT	0.409	1	N	-	-		2017-2018		N/A
	6a	LR	0.667	2	N	-	0h 1m	Add CV, incorrect undersampling			
	6b	RF	0.609	2	N	-	0h 25m				
	6c	GBT	0.684	2	N	-	12h 20m				
	7a	LR	0.4	2	N	-	0h 1m	Add CV, PROB_DEP_DEL15_ORIGIN_PRIOR_1H			
	7b	RF	0.37	2	N	-	0h 2m				
	7c	GBT	0.438	2	N	-	-				
	8	LR	0.479	5	Y	27	3h 52m	Expand CV, add gridsearch on max iter, alpha, lambda			
	9a	LR	0.445	5	Y	27	6h 30m	Further Expand CV			
	9b	LR	0.477	N/A	N/A	N/A	0h 12m	First true Test iter = 64, alpha = .2, lambda = .01 P = .671			
IV	10a	LR	0.524	N/A	N/A	256	0h 14m	iter = 64, alpha = .2, lambda = .01, threshold = .7	2015-2018		2019
	10b	RF	0.373	N/A	N/A	27	0h 24m	num_trees = 128, max_depth = 10, max_bins = 32			
	10c	GBT	0.563	N/A	N/A	3125	1h 5m	iter = 6, max_depth = 10, max_bins = 256, step_size = .2, threshold = .77 All hyperparam combos found through prior gridsearch			



Performance

Model Performance

Logistic Regression

		Predicted	
		0	1
Actual	0	0.93	0.07
	1	0.60	0.40

F0.5: 0.524

Random Forest

		Predicted	
		0	1
Actual	0	0.69	0.31
	1	0.32	0.68

F0.5: 0.373

Gradient Boosted Trees

		Predicted	
		0	1
Actual	0	0.95	0.05
	1	0.61	0.39

F0.5: 0.563

Pipeline Performance

- 80GB - 20 Cores
- Join: **5m**
- Caching: **5-10m** x checkpoint
- Total Join pipeline: **1hr**



Gap Analysis & Future Work

- Expanded Domain Knowledge
 - Aide in best feature selection/creation
- Augment with additional data sources
- Algorithmic Feature Selection/Dimensionality Reduction
 - Are we providing the model with the best information?
- Additional EDA to better inform Transformations
- Seasonality

Limitations, Challenges & Future Work

- Limitations/Challenges
 - Would like to spend more time creating/tuning engineered features
 - F0.5 2x baseline score
 - Still low after all feat eng/transformations (.563)
 - Addressing Seasonality
 - No GPU = No Neural Nets
- Future Work
 - Additional EDA to inform specific transformations
 - Interview with industry experts to gain domain knowledge
 - Address seasonality



Questions?